

missing-values-22mcb1002

March 18, 2023

1 Missing Data Handling

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

[2]: df = pd.read_csv('/content/drive/MyDrive/Data Analytics/heart.csv', na_values = ['??', '????'])
df.head(5)
```

```
[2]:    id  age     sex         cp  trestbps  chol      restecg  thalch \
0    1   63  Male  typical  angina    145.0  233.0  lv hypertrophy  150.0
1    2   67  Male  asymptomatic  160.0  286.0  lv hypertrophy  108.0
2    3   67  Male  asymptomatic  120.0  229.0  lv hypertrophy  129.0
3    4   37  Male  non-anginal  130.0  250.0      normal  187.0
4    5   41 Female  atypical  angina    130.0  204.0  lv hypertrophy  172.0

      exang  oldpeak      slope  ca
0  False       2.3  downsloping  0.0
1  True        1.5       flat    3.0
2  True        2.6       flat    2.0
3  False       3.5  downsloping  0.0
4  False       1.4  upsloping  0.0
```

```
[39]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 920 entries, 0 to 919
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   id          920 non-null    int64  
 1   age          920 non-null    int64  
 2   sex          920 non-null    object  
 3   cp           913 non-null    object  
 4   trestbps    861 non-null    float64
 5   chol         890 non-null    float64
 6   fbs          920 non-null    int64  
 7   restecg     920 non-null    object  
 8   thalch       920 non-null    int64  
 9   exercise    920 non-null    object  
 10  oldpeak      920 non-null    float64
 11  slope        920 non-null    object  

```

```
6   restecg    918 non-null    object
7   thalch     864 non-null    float64
8   exang      865 non-null    object
9   oldpeak    858 non-null    float64
10  slope       611 non-null    object
11  ca          309 non-null    float64
dtypes: float64(5), int64(2), object(5)
memory usage: 86.4+ KB
```

```
[40]: df.describe()
```

```
[40]:          id        age      trestbps       chol      thalch    oldpeak \
count  920.000000  920.000000  861.000000  890.000000  864.000000  858.000000
mean   460.500000  53.510870  132.132404  199.130337  137.538194  0.878788
std    265.725422  9.424685  19.066070  110.780810  25.940362  1.091226
min    1.000000  28.000000  0.000000  0.000000  60.000000 -2.600000
25%   230.750000  47.000000  120.000000  175.000000  120.000000  0.000000
50%   460.500000  54.000000  130.000000  223.000000  140.000000  0.500000
75%   690.250000  60.000000  140.000000  268.000000  157.000000  1.500000
max   920.000000  77.000000  200.000000  603.000000  202.000000  6.200000

           ca
count  309.000000
mean   0.676375
std    0.935653
min    0.000000
25%   0.000000
50%   0.000000
75%   1.000000
max   3.000000
```

```
[21]: df.shape
```

```
[21]: (920, 12)
```

```
[23]: df_na = df.dropna()
df_na.shape

# In the original dataset we had 920 rows containing data. But after dropping
# rows containing NaN or NULL values we are left with 302 rows.
```

```
[23]: (302, 12)
```

```
[42]: df_index = df.set_index('id')
df_index.head(5)

# Earlier we had default index starting from 0. By setting the index we had
```

```
# converted our column (id) as an index.
```

```
[42]:      age      sex          cp  trestbps  chol      restecg  thalch \
id
1   63    Male  typical angina      145.0  233.0  lv hypertrophy  150.0
2   67    Male  asymptomatic     160.0  286.0  lv hypertrophy  108.0
3   67    Male  asymptomatic     120.0  229.0  lv hypertrophy  129.0
4   37    Male  non-anginal      130.0  250.0      normal  187.0
5   41  Female  atypical angina     130.0  204.0  lv hypertrophy  172.0

      exang  oldpeak      slope      ca
id
1  False      2.3  downsloping  0.0
2  True       1.5        flat  3.0
3  True       2.6        flat  2.0
4  False      3.5  downsloping  0.0
5  False      1.4  upsloping  0.0
```

```
[45]: df_reset = df_index.reset_index('id')
df_reset.head(5)
```

```
# Resets the index to its default form.
```

```
[45]:      id  age      sex          cp  trestbps  chol      restecg  thalch \
0   1   63    Male  typical angina      145.0  233.0  lv hypertrophy  150.0
1   2   67    Male  asymptomatic     160.0  286.0  lv hypertrophy  108.0
2   3   67    Male  asymptomatic     120.0  229.0  lv hypertrophy  129.0
3   4   37    Male  non-anginal      130.0  250.0      normal  187.0
4   5   41  Female  atypical angina     130.0  204.0  lv hypertrophy  172.0

      exang  oldpeak      slope      ca
0  False      2.3  downsloping  0.0
1  True       1.5        flat  3.0
2  True       2.6        flat  2.0
3  False      3.5  downsloping  0.0
4  False      1.4  upsloping  0.0
```

```
[23]: df.isnull().any()
```

```
# This function checks if any attribute contains NaN values. If there is a
# presence of even a single NaN value it marks that attribute with boolean value
# 'True' and if there is no such value then it is marked with false. Based on
# our dataset id, age and sex are the columns with no NaN value.
```

```
[23]: id      False
age     False
sex     False
```

```
cp          True
trestbps   True
chol        True
restecg    True
thalch     True
exang      True
oldpeak    True
slope       True
ca          True
dtype: bool
```

```
[24]: df.isnull().sum()
```

```
# It counts the total NaN and NULL values present in each attribute. The oldpeak
# attribute contains the maximum number of NULL values.
```

```
[24]: id          0
age         0
sex         0
cp          7
trestbps   59
chol        30
restecg    2
thalch     56
exang      55
oldpeak    62
slope       309
ca          611
dtype: int64
```

```
[20]: df1 = df.dropna(how = 'all')
df1.shape
```

```
# It checks and drops that row in which the whole row contains NULL values.
# Based on the output there were no such rows containing all NULL values as the
# number of rows are same as original dataframe.
```

```
[20]: (920, 12)
```

```
[34]: df2 = df.dropna(axis = 1)
df2.shape
```

```
# It checks and removes that column which contains at least one missing value.
# If we go through the output we observe there are only three columns left
# that means rest all column possessed at least one missing value. Since each
# attribute plays a major role in data analysis it important to remove the
# particular data of the whole dataset and hence maintaining its integrity.
```

```
[34]: (920, 3)
```

2 Detection of missing/invalid values

```
[35]: df['cp'].unique()
```

```
# It checks all the unique values present in a particular attribute . The cp  
# attribute has four unique values and five if we count NaN value as well.
```

```
[35]: array(['typical angina', 'asymptomatic', 'non-anginal', 'atypical angina',  
           nan], dtype=object)
```

```
[36]: df['cp'].value_counts()
```

```
# It counts the frequency of unique value present in an attribute.
```

```
[36]: asymptomatic      491  
non-anginal          204  
atypical angina     172  
typical angina       46  
Name: cp, dtype: int64
```

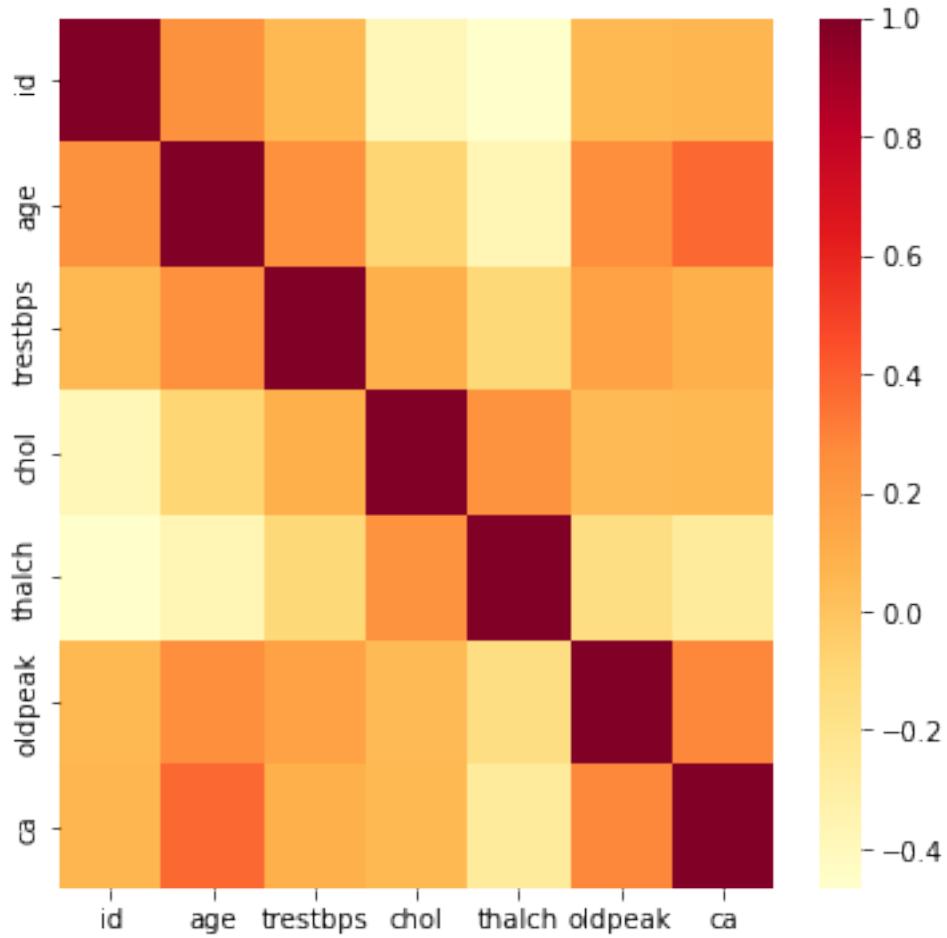
```
[5]: df['restecg'].replace('??', np.nan, inplace = True)
```

```
# It replaces ?? value with NaN value present in "restecg" column. The  
# inplace arguement validates the replace command in the original dataset rather  
# than creating a new dataset.
```

```
[52]: cormat = df.corr()  
top_cor_features = cormat.index  
plt.figure(figsize = (6,6))  
g = sns.heatmap(df[top_cor_features].corr(), cmap = "YlOrRd")
```



```
# The heatmap represents the coorelation among the attributes. The darker region  
# shows that the corresponding attribute is strongly correlated while the  
# lightest region are least coorelated. So those data which are least  
# coorelated  
# could either be removed or should be given less priority for further analysis.
```



```
[13]: missing = df[df.isnull().any(axis = 1)]
missing.head()

# It creates a datafram of all columns containing NaN or missing values. Here
# row 20, 28, 154, 166 and so on contains NaN or NULL values.
```

	id	age	sex	cp	trestbps	chol	restecg	thalch	\
20	21	64	Male	typical angina	110.0	211.0	lv hypertrophy	NaN	
28	29	43	Male		Nan	150.0	247.0	normal	171.0
154	155	64	Male		Nan	120.0	246.0	lv hypertrophy	96.0
166	167	52	Male	non-anginal	138.0	223.0	normal	169.0	
192	193	43	Male	asymptomatic	132.0	247.0	lv hypertrophy	143.0	
		exang	oldpeak	slope	ca				
20		True	1.8	flat	0.0				
28		False	1.5	upsloping	0.0				
154		True	2.2	downsloping	1.0				

```

166 False 0.0 upsloping NaN
192 True 0.1 flat NaN

```

```
[26]: df4 = df_na.astype({'thalch':int, 'chol':int})
df4.info()
```

Earlier the attribute "thalch" and "chol" was of 'float' type. After typecasting the attributes were changed to 'int' type.

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 302 entries, 0 to 748
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
---  --  
 0   id          302 non-null    int64  
 1   age         302 non-null    int64  
 2   sex         302 non-null    object  
 3   cp          302 non-null    object  
 4   trestbps   302 non-null    float64 
 5   chol        302 non-null    int64  
 6   restecg    302 non-null    object  
 7   thalch      302 non-null    int64  
 8   exang       302 non-null    object  
 9   oldpeak    302 non-null    float64 
 10  slope       302 non-null    object  
 11  ca          302 non-null    float64 
dtypes: float64(3), int64(4), object(5)
memory usage: 30.7+ KB

```

3 Filling missing values

```
[29]: df5 = df.fillna(value = 0)
df5.head()
```

It fills all the NaN values with 0. This is done during preprocessing or data cleaning in order to perform operations like regression.

```
[29]:   id  age   sex           cp  trestbps  chol           restecg  thalch \
0    1   63  Male  typical angina    145.0  233.0  lv hypertrophy  150.0
1    2   67  Male  asymptomatic  160.0  286.0  lv hypertrophy  108.0
2    3   67  Male  asymptomatic  120.0  229.0  lv hypertrophy  129.0
3    4   37  Male  non-anginal   130.0  250.0      normal  187.0
4    5   41 Female atypical angina  130.0  204.0  lv hypertrophy  172.0

      exang  oldpeak      slope  ca
0  False     2.3  downsloping  0.0
```

```

1  True      1.5      flat  3.0
2  True      2.6      flat  2.0
3 False     3.5  down sloping 0.0
4 False     1.4  up sloping 0.0

```

[30]: df5 = df.fillna(method = 'pad')
df5.head()

*# It is a forward filling method. It replaces the NaN value with the most recent
non null value present in the same column.*

[30]:

	id	age	sex	cp	trestbps	chol	restecg	thalch	\
0	1	63	Male	typical angina	145.0	233.0	lv hypertrophy	150.0	
1	2	67	Male	asymptomatic	160.0	286.0	lv hypertrophy	108.0	
2	3	67	Male	asymptomatic	120.0	229.0	lv hypertrophy	129.0	
3	4	37	Male	non-anginal	130.0	250.0	normal	187.0	
4	5	41	Female	atypical angina	130.0	204.0	lv hypertrophy	172.0	
				exang oldpeak	slope	ca			
0				False	2.3	downsloping	0.0		
1				True	1.5	flat	3.0		
2				True	2.6	flat	2.0		
3				False	3.5	downsloping	0.0		
4				False	1.4	upsloping	0.0		

[38]: df5 = df.fillna(method = 'bfill')
df5.head()

*# It is a backward filling method. It replaces the NaN value with the next
non null value present in the same column.*

[38]:

	id	age	sex	cp	trestbps	chol	restecg	thalch	\
0	1	63	Male	typical angina	145.0	233.0	lv hypertrophy	150.0	
1	2	67	Male	asymptomatic	160.0	286.0	lv hypertrophy	108.0	
2	3	67	Male	asymptomatic	120.0	229.0	lv hypertrophy	129.0	
3	4	37	Male	non-anginal	130.0	250.0	normal	187.0	
4	5	41	Female	atypical angina	130.0	204.0	lv hypertrophy	172.0	
				exang oldpeak	slope	ca			
0				False	2.3	downsloping	0.0		
1				True	1.5	flat	3.0		
2				True	2.6	flat	2.0		
3				False	3.5	downsloping	0.0		
4				False	1.4	upsloping	0.0		

[54]: df.interpolate()
df.head()

```
# It replaces the missing values using linear interpolation. This is done by
# observing the neighboring value present in that column.
```

```
[54]:   id  age    sex      cp  trestbps  chol      restecg  thalch \
0    1    63  Male  typical angina    145.0  233.0  lv hypertrophy  150.0
1    2    67  Male  asymptomatic  160.0  286.0  lv hypertrophy  108.0
2    3    67  Male  asymptomatic  120.0  229.0  lv hypertrophy  129.0
3    4    37  Male  non-anginal  130.0  250.0      normal  187.0
4    5    41  Female atypical angina  130.0  204.0  lv hypertrophy  172.0

      exang  oldpeak      slope  ca
0  False      2.3  downsloping  0.0
1  True       1.5        flat  3.0
2  True       2.6        flat  2.0
3  False      3.5  downsloping  0.0
4  False      1.4  upsloping  0.0
```

```
[36]: df.interpolate().count()
```

```
# After the interpolate() method fills the missing or null values in df, the
# count() method is used to count the number of non-null values in each column
# of the resulting DataFrame.
```

```
[36]: id      920
age     920
sex     920
cp      913
trestbps 920
chol    920
restecg  918
thalch   920
exang    865
oldpeak  920
slope    611
ca      920
dtype: int64
```

```
[37]: df.interpolate().plot()
```

```
# After the interpolate() method fills the missing or null values in df, the
# plot() method is used to plot the non-null values in each column
# of the resulting DataFrame.
```

```
[37]: <Axes: >
```

