

22mcb1002

April 7, 2023

```
[65]: import pandas as pd
import numpy as np
from sklearn.preprocessing import MaxAbsScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import RobustScaler

[2]: df = pd.read_csv('/content/drive/MyDrive/Data Analytics/heart.csv')

[3]: df.columns

[3]: Index(['id', 'age', 'sex', 'cp', 'trestbps', 'chol', 'restecg', 'thalch',
       'exang', 'oldpeak', 'slope', 'ca'],
       dtype='object')

[4]: df.head()

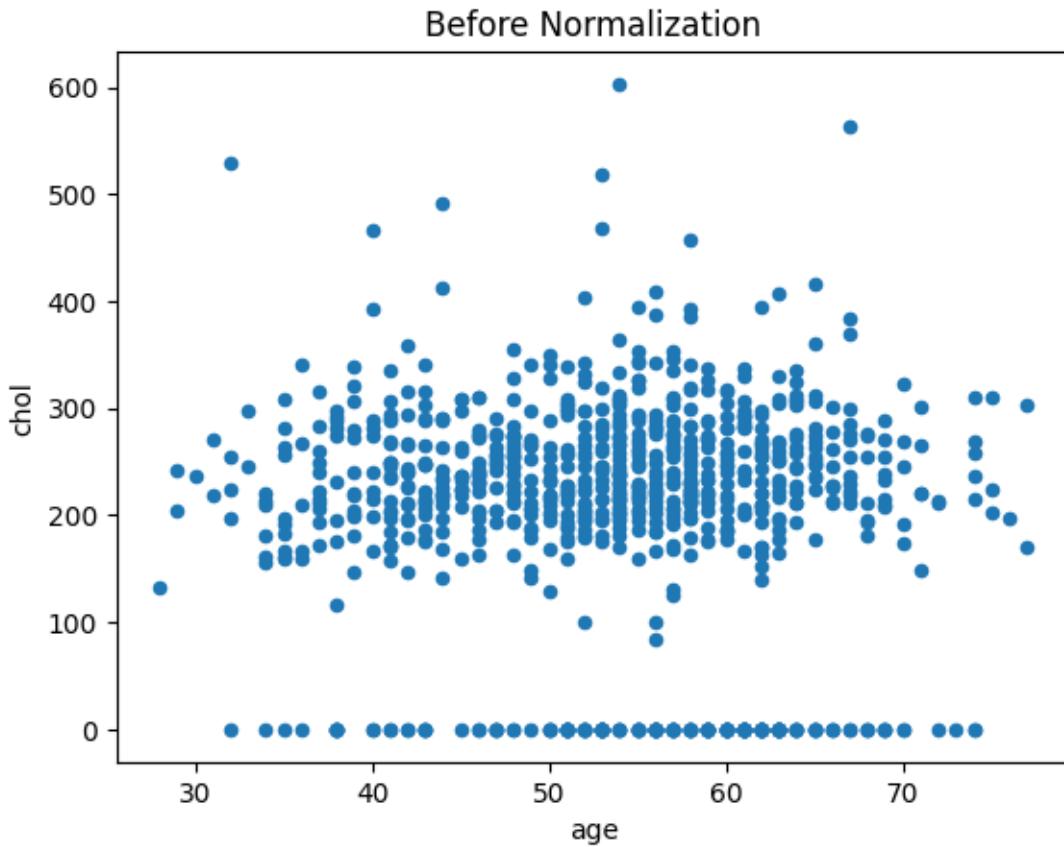
[4]:   id  age     sex         cp trestbps    chol      restecg thalch \
0    1   63    Male  typical angina    145.0  233.0    lv hypertrophy   150.0
1    2   67    Male asymptomatic    160.0  286.0    lv hypertrophy   108.0
2    3   67    Male asymptomatic    120.0  229.0    lv hypertrophy   129.0
3    4   37    Male non-anginal    130.0  250.0      normal    187.0
4    5   41  Female atypical angina    130.0  204.0    lv hypertrophy   172.0

      exang  oldpeak      slope  ca
0  False       2.3  downslloping  0.0
1  True        1.5          flat  3.0
2  True        2.6          flat  2.0
3  False       3.5  downslloping  0.0
4  False       1.4      upsloping  0.0

[39]: df.plot.scatter(x = 'age', y = 'chol', title = 'Before Normalization')

/usr/local/lib/python3.9/dist-packages/pandas/plotting/_matplotlib/core.py:1114:
UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap' will
be ignored
scatter = ax.scatter(
```

```
[39]: <Axes: title={'center': 'Before Normalization'}, xlabel='age', ylabel='chol'>
```



```
[7]: df.columns[:4]
```

```
[7]: Index(['id', 'age', 'sex', 'cp'], dtype='object')
```

```
[9]: data_numeric = df.select_dtypes(include = 'number')
data_numeric.columns
```

```
[9]: Index(['id', 'age', 'trestbps', 'chol', 'thalch', 'oldpeak', 'ca'],
      dtype='object')
```

1 Maximum absolute scaling

Without using predefined function

```
[49]: df1 = df[['age', 'chol']]
def max_abs_scaling(df1):
    df_scaled = df1.copy()
```

```
for col in df_scaled.columns[:2]:  
    df_scaled[col] = df_scaled[col]/df_scaled[col].abs().max()  
return df_scaled  
df_cars_scaled = max_abs_scaling(df1)  
df_cars_scaled.head()
```

[49]:

	age	chol
0	0.818182	0.386401
1	0.870130	0.474295
2	0.870130	0.379768
3	0.480519	0.414594
4	0.532468	0.338308

[28]:

```
abs_scaler = MaxAbsScaler()  
df1 = df[['age', 'chol']]  
df1.columns
```

[28]:

```
Index(['age', 'chol'], dtype='object')
```

Calculating maximum absolute value

[30]:

```
abs_scaler.fit(df1)  
MaxAbsScaler(copy = True)
```

[30]:

```
MaxAbsScaler()
```

[31]:

```
abs_scaler.max_abs_
```

[31]:

```
array([ 77., 603.])
```

Transforming data

[33]:

```
scaled_data = abs_scaler.transform(df1)  
scaled_data.dtype
```

[33]:

```
dtype('float64')
```

Using predefined function

[42]:

```
df_scaled_data = pd.DataFrame(scaled_data, columns = df1.columns)  
df_scaled_data.head()
```

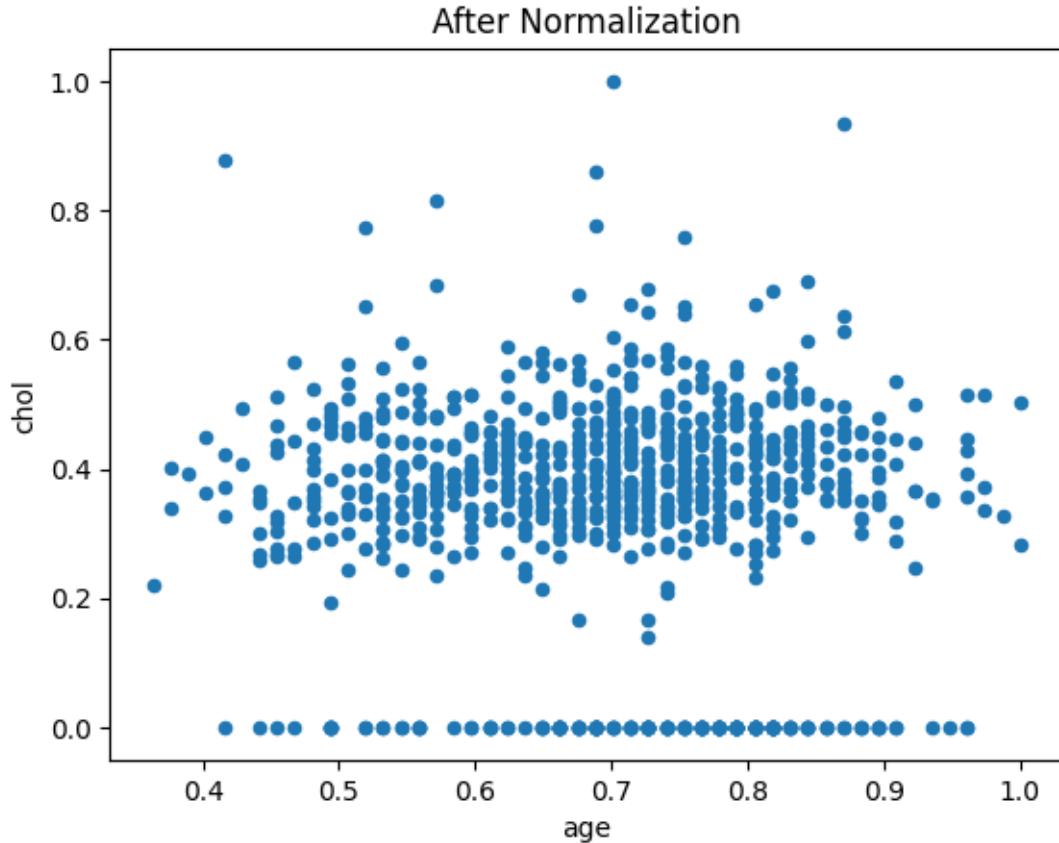
[42]:

	age	chol
0	0.818182	0.386401
1	0.870130	0.474295
2	0.870130	0.379768
3	0.480519	0.414594
4	0.532468	0.338308

```
[43]: df_scaled.plot.scatter(x = 'age', y = 'chol', title = 'After  
Normalization')
```

```
/usr/local/lib/python3.9/dist-packages/pandas/plotting/_matplotlib/core.py:1114:  
UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap' will  
be ignored  
    scatter = ax.scatter(
```

```
[43]: <Axes: title={'center': 'After Normalization'}, xlabel='age', ylabel='chol'>
```



2 Min-Max Normalisation

Without using predefined function

```
[45]: def min_max_scaling(df):  
    df_scaled = df.copy()  
    for col in df_scaled.columns:  
        df_scaled[col] = (df_scaled[col]-df_scaled[col].min())/(df_scaled[col].  
        max()-df_scaled[col].min())
```

```
    return df_scaled
df_cars_scaled = min_max_scaling(df1)
df_cars_scaled.head()
```

[45]:

	age	chol
0	0.714286	0.386401
1	0.795918	0.474295
2	0.795918	0.379768
3	0.183673	0.414594
4	0.265306	0.338308

Using predefined function

[48]:

```
MMscaler = MinMaxScaler()
df_norm = pd.DataFrame(MMscaler.fit_transform(df1), columns = df1.columns)
df_norm.head()
```

[48]:

	age	chol
0	0.714286	0.386401
1	0.795918	0.474295
2	0.795918	0.379768
3	0.183673	0.414594
4	0.265306	0.338308

[51]:

```
MMscaler.data_min_
#MMscaler.data_max_
```

[51]:

```
array([28.,  0.])
```

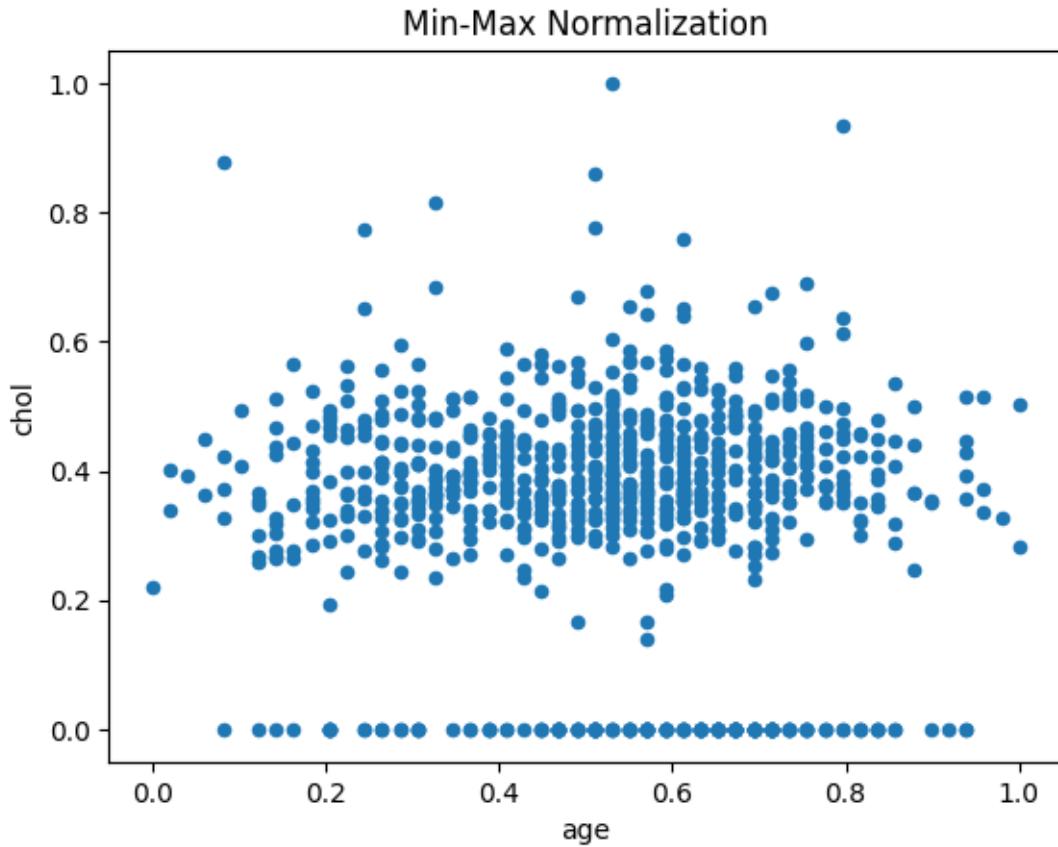
[52]:

```
df_norm.plot.scatter(x = 'age', y = 'chol', title = 'Min-Max Normalization')
```

```
/usr/local/lib/python3.9/dist-packages/pandas/plotting/_matplotlib/core.py:1114:
UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap' will
be ignored
    scatter = ax.scatter(
```

[52]:

```
<Axes: title={'center': 'Min-Max Normalization'}, xlabel='age', ylabel='chol'>
```



3 Z Score method

Without using predefined function

```
[53]: def z_norm(df):
    df_scaled=df.copy()
    for col in df_scaled.columns:
        df_scaled[col]=(df_scaled[col]-df_scaled[col].mean())/(df_scaled[col].std())
    return df_scaled
df_scaled_cars = z_norm(df1)
df_scaled_cars.head()
```

```
[53]:      age      chol
0  1.006838  0.305736
1  1.431255  0.784158
2  1.431255  0.269628
3 -1.751875  0.459192
4 -1.327458  0.043958
```

Using predefined function

```
[55]: stdscaler = StandardScaler()
stdscaler_data = stdscaler.fit_transform(df1)
stdscaler_df = pd.DataFrame(stdscaler_data, columns = df1.columns)
stdscaler_df.head()
```

```
[55]:      age      chol
0  1.007386  0.305908
1  1.432034  0.784599
2  1.432034  0.269780
3 -1.752828  0.459450
4 -1.328180  0.043982
```

```
[56]: stdscaler.scale_
```

```
[56]: array([ 9.41956171, 110.71855645])
```

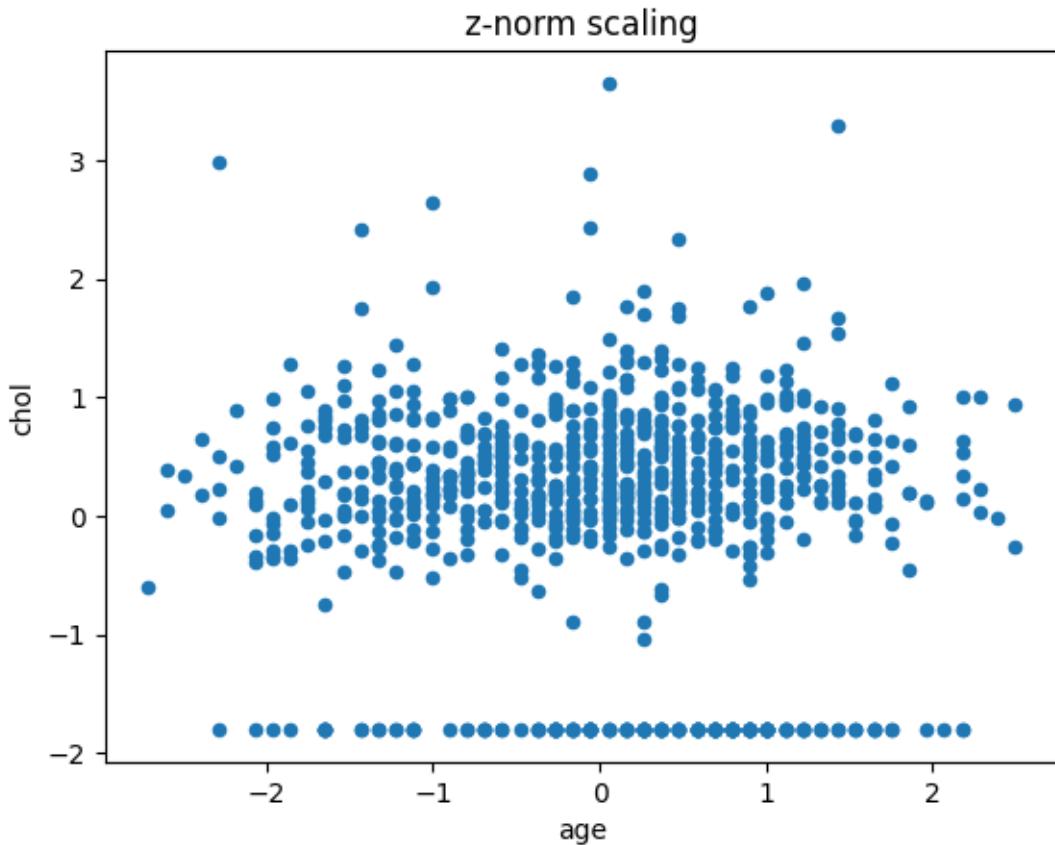
```
[57]: stdscaler.mean_
```

```
[57]: array([ 53.51086957, 199.13033708])
```

```
[58]: stdscaler_df.plot.scatter(x = 'age', y = 'chol', title = "z-norm scaling")
```

```
/usr/local/lib/python3.9/dist-packages/pandas/plotting/_matplotlib/core.py:1114:
UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap' will
be ignored
    scatter = ax.scatter(
```

```
[58]: <Axes: title={'center': 'z-norm scaling'}, xlabel='age', ylabel='chol'>
```



4 Robust Scaling

Without using predefined function

```
[59]: def robust_scaling(df):
    df_scaled = df.copy()
    for col in df_scaled.columns:
        df_scaled[col] = (df_scaled[col]-df_scaled[col].median())/(df_scaled[col].quantile(0.75)-df_scaled[col].quantile(0.25))
    return df_scaled
df_scaled=robust_scaling(df1)
df_scaled.head()
```

	age	chol
0	0.692308	0.107527
1	1.000000	0.677419
2	1.000000	0.064516
3	-1.307692	0.290323
4	-1.000000	-0.204301

Using predefined function

```
[62]: rob_scale = RobustScaler()
rob_scaledata = rob_scale.fit_transform(df1)
rob_scaled_df = pd.DataFrame(rob_scaledata, columns = df1.columns)
rob_scaled_df.head()
```

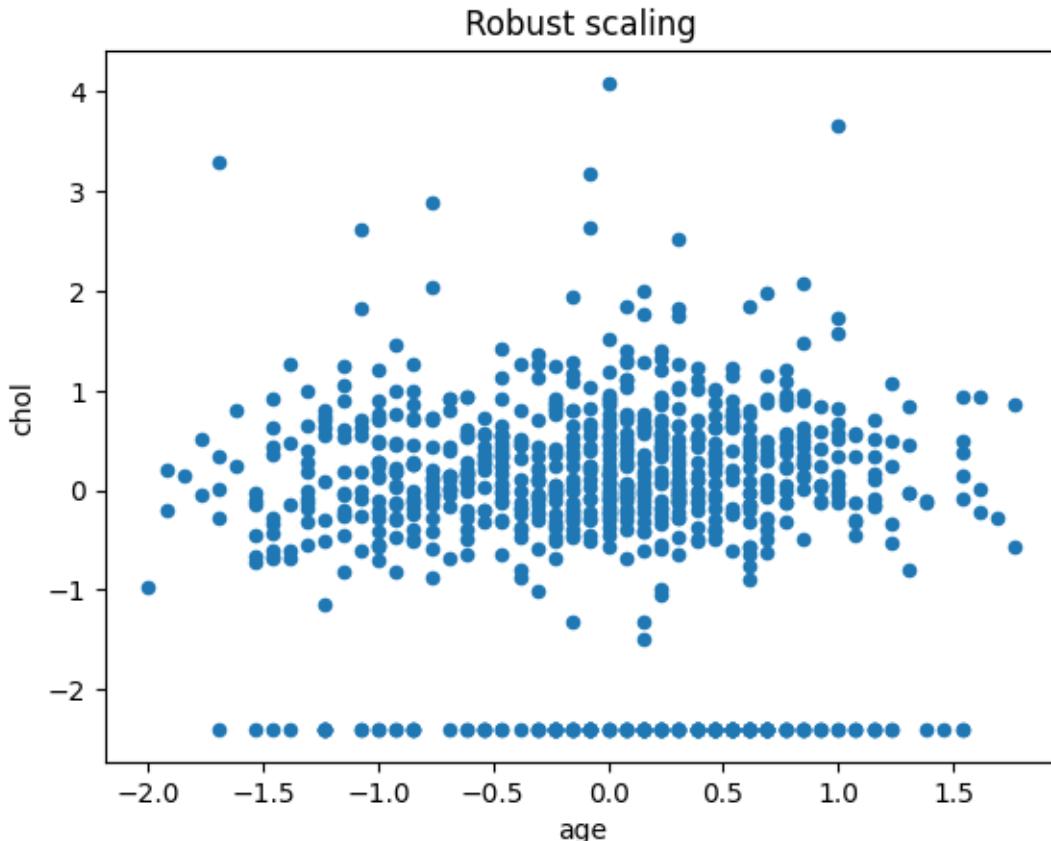
```
[62]:      age      chol
0  0.692308  0.107527
1  1.000000  0.677419
2  1.000000  0.064516
3 -1.307692  0.290323
4 -1.000000 -0.204301
```

```
[63]: rob_scaled_df.plot.scatter(x = 'age', y = 'chol', title = 'Robust scaling')
```

```
/usr/local/lib/python3.9/dist-packages/pandas/plotting/_matplotlib/core.py:1114:
UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap' will
be ignored
```

```
    scatter = ax.scatter(
```

```
[63]: <Axes: title={'center': 'Robust scaling'}, xlabel='age', ylabel='chol'>
```



5 Binning

Binning by distance - use `cut()`

```
[64]: min_age = df1['age'].min()  
       max_age = df1['age'].max()  
       print(min_age)  
       print(max_age)
```

28
77

```
[66]: # Returns equally spaced values  
       np.linspace(1,10, 4)
```

```
[66]: array([ 1.,  4.,  7., 10.])
```

```
[67]: bins = np.linspace(min_age, max_age, 4)  
       bins
```

```
[67]: array([28.           , 44.33333333, 60.66666667, 77.           ])
```

```
[69]: labels =['young','middle','senior']  
       df1.columns
```

```
[69]: Index(['age', 'chol'], dtype='object')
```

```
[70]: df1['Age_categ'] = pd.cut(df1['age'], bins = bins, labels = labels,  
                                include_lowest = True)  
       df1.columns
```

```
<ipython-input-70-017515a25297>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
`df1['Age_categ'] = pd.cut(df1['age'], bins = bins, labels = labels,
include_lowest = True)`

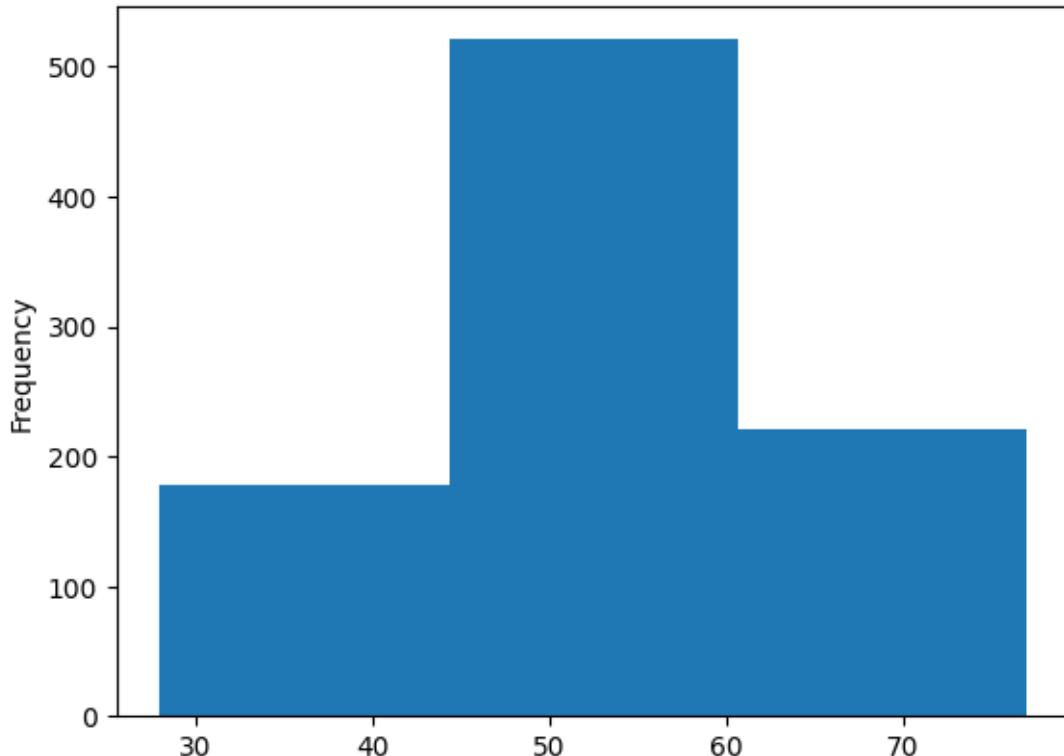
```
[70]: Index(['age', 'chol', 'Age_categ'], dtype='object')
```

```
[71]: df1.head()
```

```
[71]:    age    chol Age_categ
0    63   233.0    senior
1    67   286.0    senior
2    67   229.0    senior
3    37   250.0    young
4    41   204.0    young
```

```
[72]: df1['age'].plot.hist(bins = 3)
```

```
[72]: <Axes: ylabel='Frequency'>
```



Binning by frequency - use `qcut()`

```
[73]: df1['Age_categ_freq'] = pd.qcut(df1['age'], q = 3, labels = labels, precision = 1)
df1.columns
```

```
<ipython-input-73-0bf73af84618>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df1['Age_categ_freq'] = pd.qcut(df1['age'], q = 3, labels = labels, precision = 1)
```

[73]: Index(['age', 'chol', 'Age_categ', 'Age_categ_freq'], dtype='object')

[74]: df1.head()

```
[74]:   age    chol Age_categ Age_categ_freq
 0    63  233.0    senior      senior
 1    67  286.0    senior      senior
 2    67  229.0    senior      senior
 3    37  250.0    young       young
 4    41  204.0    young       young
```